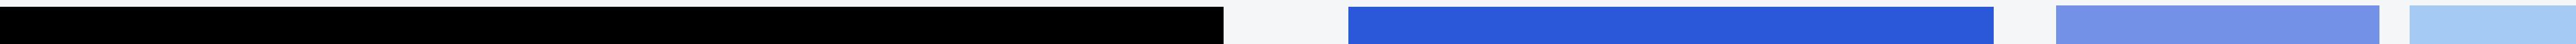


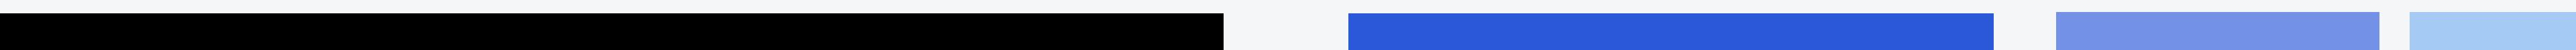
Lecture-2

Markov Decision

Process




Agenda



01. Markov Process
02. Reward and Return
03. Value Functions
04. Bellman Equations

01.

Markov Process



Markov Property

- The future is independent of the past given the present.

A state S_t is *Markov* if and only if

$$\mathbb{P}[S_{t+1} \mid S_t] = \mathbb{P}[S_{t+1} \mid S_1, \dots, S_t]$$

- The state captures all relevant information from the history
- Once the state is known, the history may be thrown away
- The state is a sufficient statistic of the future

Markov Chain

- A Markov chain is a memoryless random process, i.e. a sequence of random states S_1, S_2, \dots with the Markov property.

Definition

A *Markov Process* (or *Markov Chain*) is a tuple $\langle \mathcal{S}, \mathcal{P} \rangle$

- \mathcal{S} is a (finite) set of states
- \mathcal{P} is a state transition probability matrix,
$$\mathcal{P}_{ss'} = \mathbb{P}[S_{t+1} = s' \mid S_t = s]$$

State Transition Matrix


- State Transition Probability

$$\mathcal{P}_{ss'} = \mathbb{P} [S_{t+1} = s' \mid S_t = s]$$

- State Transition Matrix


$$\mathcal{P} = \begin{matrix} & & \text{to} \\ & & \left[\begin{matrix} \mathcal{P}_{11} & \dots & \mathcal{P}_{1n} \\ \vdots & & \vdots \\ \mathcal{P}_{n1} & \dots & \mathcal{P}_{nn} \end{matrix} \right] \\ \text{from} & & \end{matrix}$$

State Transition Matrix



$$\mathcal{P} = \begin{bmatrix} C1 & C2 & C3 & Pass & Pub & FB & Sleep \\ C1 & 0.5 & 0.8 & 0.6 & 0.4 & 0.9 & 1.0 \\ C2 & 0.2 & 0.4 & 0.6 & 0.4 & 0.5 & 0.2 \\ C3 & 0.1 & 0.4 & 0.4 & 0.4 & 0.4 & 0.1 \\ Pass & 0.2 & 0.4 & 0.4 & 0.4 & 0.4 & 0.2 \\ Pub & 0.1 & 0.1 & 0.1 & 0.1 & 0.1 & 0.1 \\ FB & 0.1 & 0.1 & 0.1 & 0.1 & 0.1 & 0.1 \\ Sleep & 0.1 & 0.1 & 0.1 & 0.1 & 0.1 & 0.1 \end{bmatrix}$$

Episode Sampling



- C1 C2 C3 Pass Sleep
- C1 FB FB C1 C2 Sleep
- C1 C2 C3 Pub C2 C3 Pass Sleep
- C1 FB FB C1 C2 C3 Pub C1 FB FB
FB C1 C2 C3 Pub C2 Sleep

Markov Reward Process


- A Markov reward process is a Markov chain with values.

Definition

A *Markov Reward Process* is a tuple $\langle \mathcal{S}, \mathcal{P}, \mathcal{R}, \gamma \rangle$

- \mathcal{S} is a finite set of states
- \mathcal{P} is a state transition probability matrix,
 $\mathcal{P}_{ss'} = \mathbb{P}[S_{t+1} = s' \mid S_t = s]$
- \mathcal{R} is a reward function, $\mathcal{R}_s = \mathbb{E}[R_{t+1} \mid S_t = s]$
- γ is a discount factor, $\gamma \in [0, 1]$

Markov Reward Process



Markov Decision Process

- A Markov decision process (MDP) is a Markov reward process with decisions.
- It is an environment in which all states are Markov.

Definition

A *Markov Decision Process* is a tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$


- \mathcal{S} is a finite set of states
- \mathcal{A} is a finite set of actions
- \mathcal{P} is a state transition probability matrix,
 $\mathcal{P}_{ss'}^{\textcolor{red}{a}} = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = \textcolor{red}{a}]$
- \mathcal{R} is a reward function, $\mathcal{R}_s^{\textcolor{red}{a}} = \mathbb{E}[R_{t+1} | S_t = s, A_t = \textcolor{red}{a}]$
- γ is a discount factor $\gamma \in [0, 1]$.

Markov Decision Process

$$p(s', r | s, a) \doteq \Pr\{S_t = s', R_t = r \mid S_{t-1} = s, A_{t-1} = a\},$$


$$\sum_{s' \in \mathcal{S}} \sum_{r \in \mathcal{R}} p(s', r | s, a) = 1, \text{ for all } s \in \mathcal{S}, a \in \mathcal{A}(s).$$

Markov Decision Process



$S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, R_3, \dots$

Markov Decision Process




$$\mathcal{S} = \{\text{high}, \text{low}\}.$$

$$\mathcal{A}(\text{high}) = \{\text{search}, \text{wait}\}$$


$$\mathcal{A}(\text{low}) = \{\text{search}, \text{wait}, \text{recharge}\}.$$

Markov Decision Process



| s | a | s' | $p(s' s, a)$ | $r(s, a, s')$ |
|------|----------|------|----------------|---------------------|
| high | search | high | α | r_{search} |
| high | search | low | $1 - \alpha$ | r_{search} |
| low | search | high | $1 - \beta$ | -3 |
| low | search | low | β | r_{search} |
| high | wait | high | 1 | r_{wait} |
| high | wait | low | 0 | - |
| low | wait | high | 0 | - |
| low | wait | low | 1 | r_{wait} |
| low | recharge | high | 1 | 0 |
| low | recharge | low | 0 | - |

Task



Give a table analogous to previous table, but for $p(s', r | s, a)$

It should have columns


$$\begin{matrix} s, a, s', r, \\ p(s', r | s, a) \end{matrix}$$

and a row for every 4-tuple for

which $p(s', r | s, a) > 0$.

02.

Reward and Return




Return

- The agent's goal is to maximize the cumulative reward in the long run.
- Cumulative reward is called Return.
- Return, denoted G_t , is defined as some specific function of the reward sequence.



Episodic Tasks



- C1 C2 C3 Pass Sleep
- C1 FB FB C1 C2 Sleep
- C1 C2 C3 Pub C2 C3 Pass Sleep
- C1 FB FB C1 C2 C3 Pub C1 FB FB
FB C1 C2 C3 Pub C2 Sleep

Episodic Tasks

- Episodic tasks: interaction breaks naturally into episodes, e.g., plays of a game, trips through a maze.

$$G_t \doteq R_{t+1} + R_{t+2} + R_{t+3} + \cdots + R_T,$$

- where T is a final time step at which a terminal state is reached, ending an episode.

Continuing Tasks

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1},$$

$$\begin{aligned} G_t &\doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \cdots \\ &= R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \gamma^2 R_{t+4} + \cdots) \\ &= R_{t+1} + \gamma G_{t+1} \end{aligned}$$


Discount Factor

- The *discount* $\gamma \in [0, 1]$ is the present value of future rewards
- The value of receiving reward R after $k + 1$ time-steps is $\gamma^k R$.
- This values immediate reward above delayed reward.
 - γ close to 0 leads to "myopic" evaluation
 - γ close to 1 leads to "far-sighted" evaluation

Task

In a Markov decision process, a large discount factor γ means that short term rewards are much more influential than long term rewards.

- True
- False



Task

Exercise 3.8 Suppose $\gamma = 0.5$ and the following sequence of rewards is received $R_1 = -1$, $R_2 = 2$, $R_3 = 6$, $R_4 = 3$, and $R_5 = 2$, with $T = 5$. What are G_0, G_1, \dots, G_5 ? Hint: Work backwards. \square

Task

Exercise 3.9 Suppose $\gamma = 0.9$ and the reward sequence is $R_1 = 2$ followed by an infinite sequence of 7s. What are G_1 and G_0 ? \square

03.

Value Function

Value Function

- Functions of states (or of state–action pairs) that estimate how good it is for the agent to be in a given state (or how good it is to perform a given action in a given state).
 - The notion of “how good” is defined in terms of expected return.
 - Value functions are defined with respect to particular policies.
- 

Value Function

State Value


$$v_{\pi}(s) \doteq \mathbb{E}_{\pi}[G_t \mid S_t = s] = \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s \right], \text{ for all } s \in \mathcal{S},$$

Action Value

$$q_{\pi}(s, a) \doteq \mathbb{E}_{\pi}[G_t \mid S_t = s, A_t = a] = \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a \right].$$

04.

Bellman Equation




The decorative horizontal bars consist of four distinct colored segments: a thick black bar on the left, a thick blue bar in the center, a thinner light blue bar to its right, and a very thin light blue bar on the far right.

Bellman Equation



$$\begin{aligned} v_\pi(s) &\doteq \mathbb{E}_\pi[G_t \mid S_t = s] \\ &= \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid S_t = s] \\ &= \sum_a \pi(a|s) \sum_{s'} \sum_r p(s', r | s, a) \left[r + \gamma \mathbb{E}_\pi[G_{t+1} | S_{t+1} = s'] \right] \\ &= \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) \left[r + \gamma v_\pi(s') \right], \quad \text{for all } s \in \mathcal{S}, \end{aligned}$$

Backup Diagram



Backup diagram for v_π

Task



- A gridworld representation of a simple finite MDP.
- The cells of the grid correspond to the states of the environment.
- At each cell, four actions are possible: north, south, east, and west.

Task

- Actions that would take the agent off the grid leave its location unchanged but also result in a reward of -1 .
- Other actions result in a reward of 0 , except those that move the agent out of the special states A and B.
- From state A, all four actions yield a reward of $+10$ and take the agent to A' .
From state B, all actions yield a reward of $+5$ and take the agent to B' .

Task

Exercise 3.14 The Bellman equation (3.14) must hold for each state for the value function v_π shown in Figure 3.2 (right) of Example 3.5. Show numerically that this equation holds for the center state, valued at +0.7, with respect to its four neighboring states, valued at +2.3, +0.4, -0.4, and +0.7. (These numbers are accurate only to one decimal place.) \square

| | | | | |
|------|------|------|------|------|
| 3.3 | 8.8 | 4.4 | 5.3 | 1.5 |
| 1.5 | 3.0 | 2.3 | 1.9 | 0.5 |
| 0.1 | 0.7 | 0.7 | 0.4 | -0.4 |
| -1.0 | -0.4 | -0.4 | -0.6 | -1.2 |
| -1.9 | -1.3 | -1.2 | -1.4 | -2.0 |

Optimal Policy and Bellman Equation


$$v_*(s) \doteq \max_{\pi} v_{\pi}(s),$$

$$\begin{aligned} v_*(s) &= \max_{a \in \mathcal{A}(s)} q_{\pi_*}(s, a) \\ &= \max_a \mathbb{E}_{\pi_*}[G_t \mid S_t = s, A_t = a] \\ &= \max_a \mathbb{E}_{\pi_*}[R_{t+1} + \gamma G_{t+1} \mid S_t = s, A_t = a] \\ &= \max_a \mathbb{E}[R_{t+1} + \gamma v_*(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_*(s')]. \end{aligned}$$

$$q_*(s, a) = \mathbb{E}[R_{t+1} + \gamma v_*(S_{t+1}) \mid S_t = s, A_t = a].$$

Task

Give the Bellman optimality equations for this recycling robot.



**Thank
You!**

